

ThinLTO: Scalable and Incremental LTO (CGO'17)

reading event

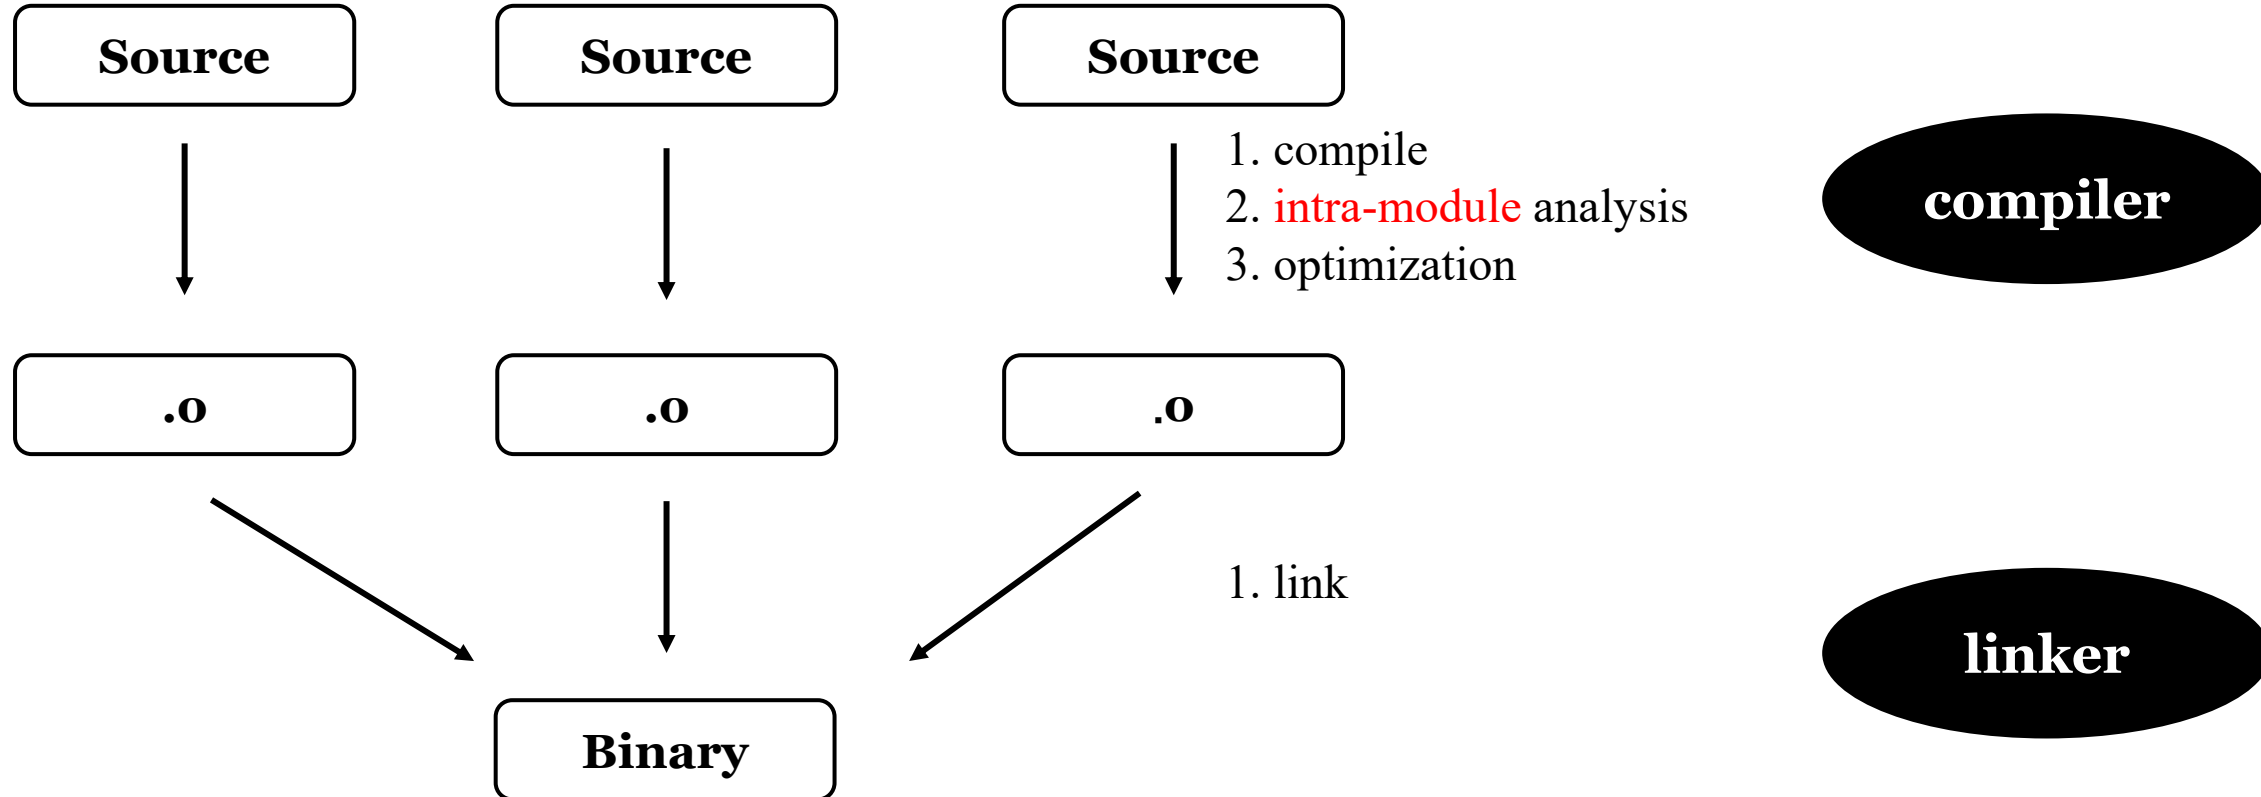
Nov. 4, 2021. Bowen Zhang



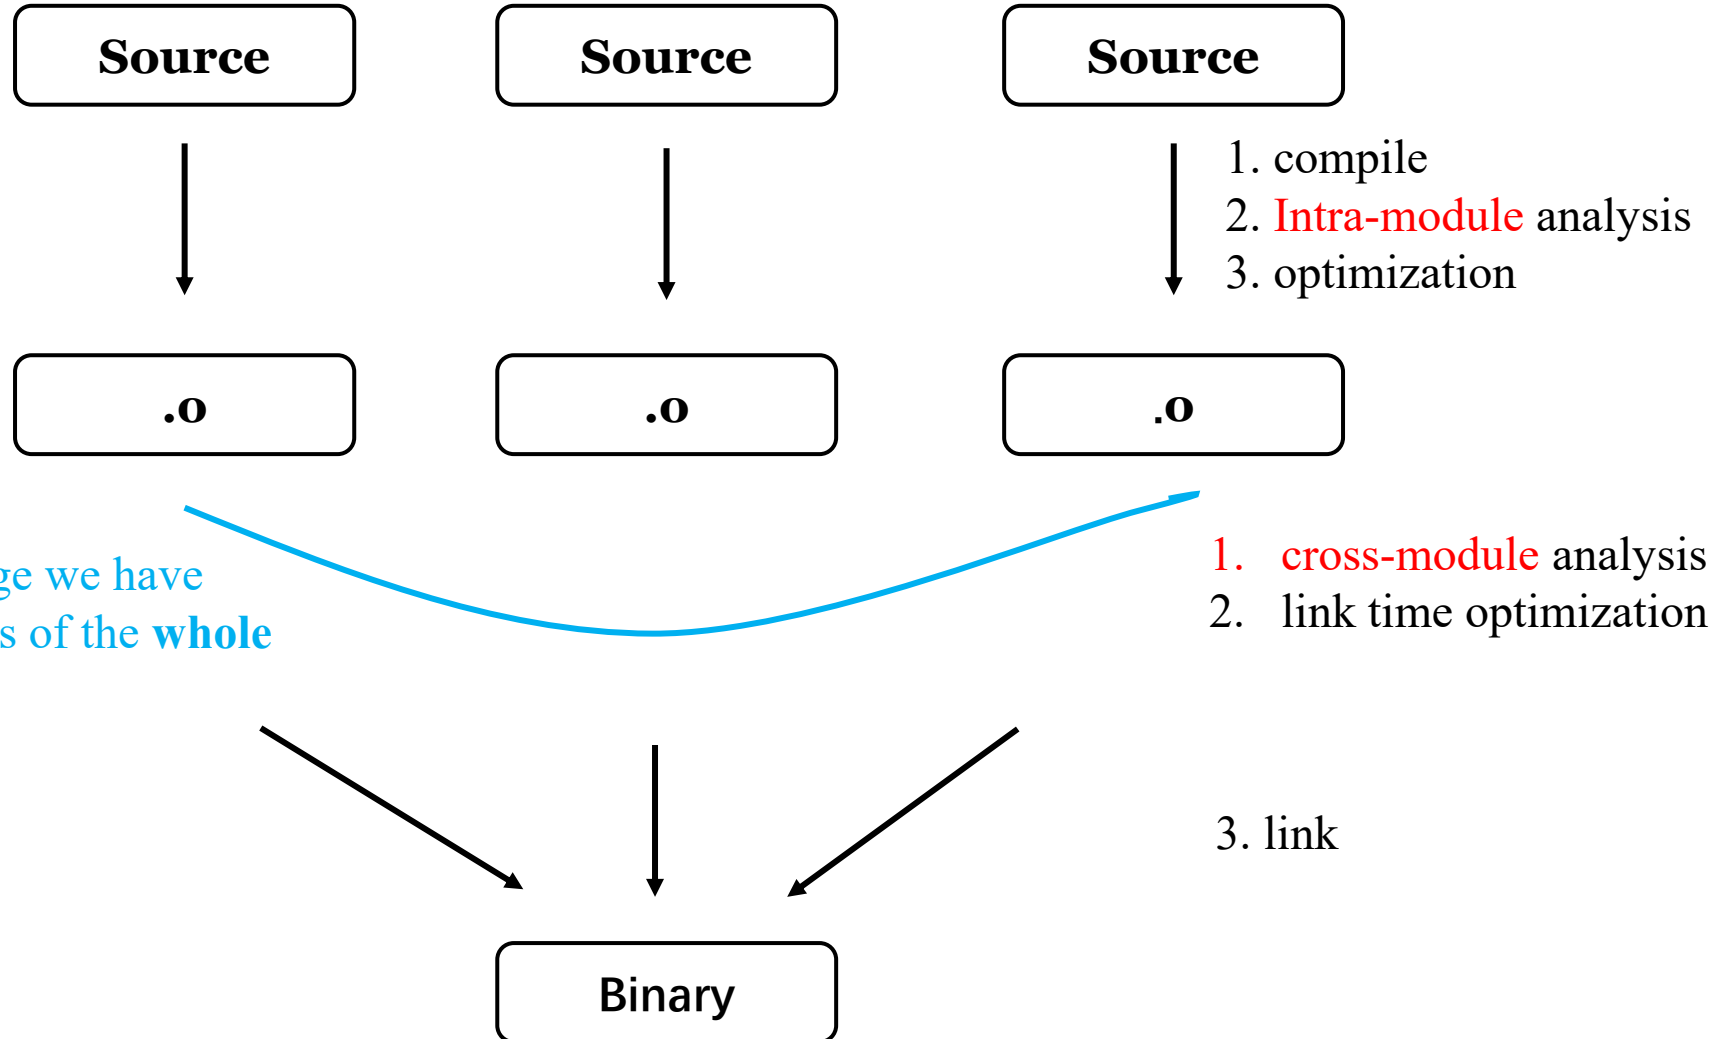
Contents

- Link Time Optimization(LTO)
 - What is LTO?
 - What benefits?
- Why LTO couldn't be enabled by default?
- How ThinLTO overcomes this...
- Discussion

Traditional Build



Build with LTO



compiler

linker

Benefits

Smaller Binary Size

- Main contribution
 - Dead Code Elimination
 - Constant Propagation

Better Performance

- 10% improvement is common.
- Main contribution: Inline
 - Usually performed with PGO(profile-guided optimization)

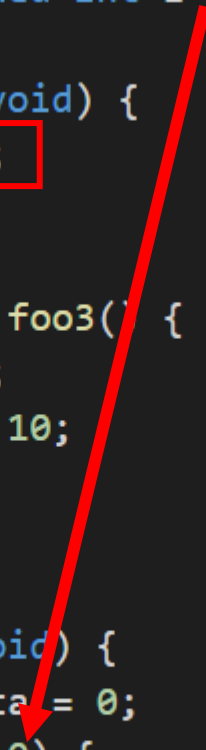
Binary Size

Can we reduce the code size?
-very hard

We don't know whether `foo2` is called at some other places, so we can't eliminate this `if-block`.

A.c

```
1  void foo4(void);
2
3  static signed int i=0;
4
5  void foo2(void) {
6      i = -1;
7  }
8
9  static int foo3() {
10     foo4();
11     return 10;
12 }
13
14
15 int foo1(void) {
16     int data = 0;
17     if(i < 0) {
18         data = foo3();
19     }
20     data = data + 42;
21     return data;
22 }
```



Binary Size

Can we reduce the code size now?

-yes

main.c

```
1  #include <stdio.h>
2
3  int foo1(void);
4
5  void foo4(void) {
6      printf("Hi\n");
7  }
8
9
10 int main() {
11     return foo1();
12 }
```

foo4 never used

(don't need to
link printf() into
the binary too!)

A.c

```
1  void foo4(void);
2
3  static signed int i=0;
4
5  void foo2(void) {
6      i = -1;
7  }
8
9  static int foo3() {
10     foo4();
11     return 10;
12 }
13
14
15 int foo1(void) {
16     int data = 0;
17     if(i < 0) {
18         data = foo3();
19     }
20     data = data + 42;
21     return data;
22 }
```

foo2 never used

foo3 never used

$i < 0$ always false

data = 42

Performance

Do we reach the end?

-no

main.c

```
1  #include <stdio.h>
2
3  int foo1(void);
4
5  void foo4(void) {
6      printf("Hi\n");
7  }
8
9
10 int main() {
11     return 42;
12 }
```

Inline

A.c

```
1  void foo4(void);
2
3  static signed int i=0;
4
5  void foo2(void) {
6      i = -1;
7  }
8
9  static int foo3() {
10     foo4();
11     return 10;
12 }
13
14
15 int foo1(void) {
16     int data = 0;
17     if(i < 0) {
18         data = foo3();
19     }
20     data = data + 42;
21     return data;
22 }
```

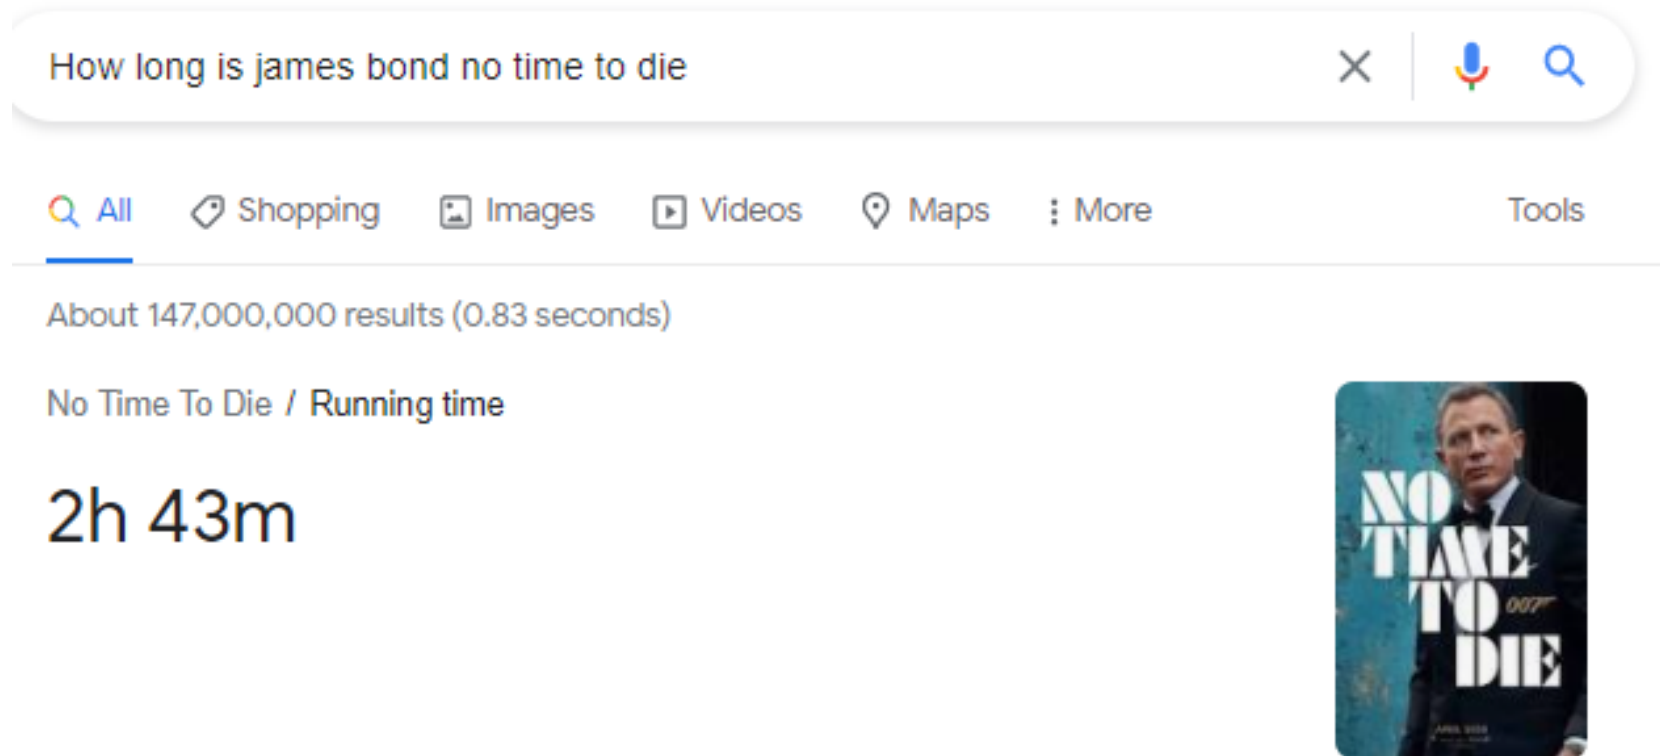
try compile this example with LTO yourself

clang -c main.c A.c -flto

clang -o main main.o A.o -flto

Can't enable LTO by default

GCC's LTO(-g0)



M's LTO can't finish this task.

which has 2X call graph nodes, and 5X call graph edges, compared with Chromium

ThinLTO: Why?
scalable & low memory lean

GCC's LTO(-g0)
ThinLTO



clang

55 s

2 GB

5 s

0.13 GB



chromium

4 mins

8 GB

30 s

1 GB

G Ad Delivery

**Google's ad recommendation application
which has 2X call graph nodes, and 3X call
graph edges, compared with Chromium*

3 hours

>25 GB

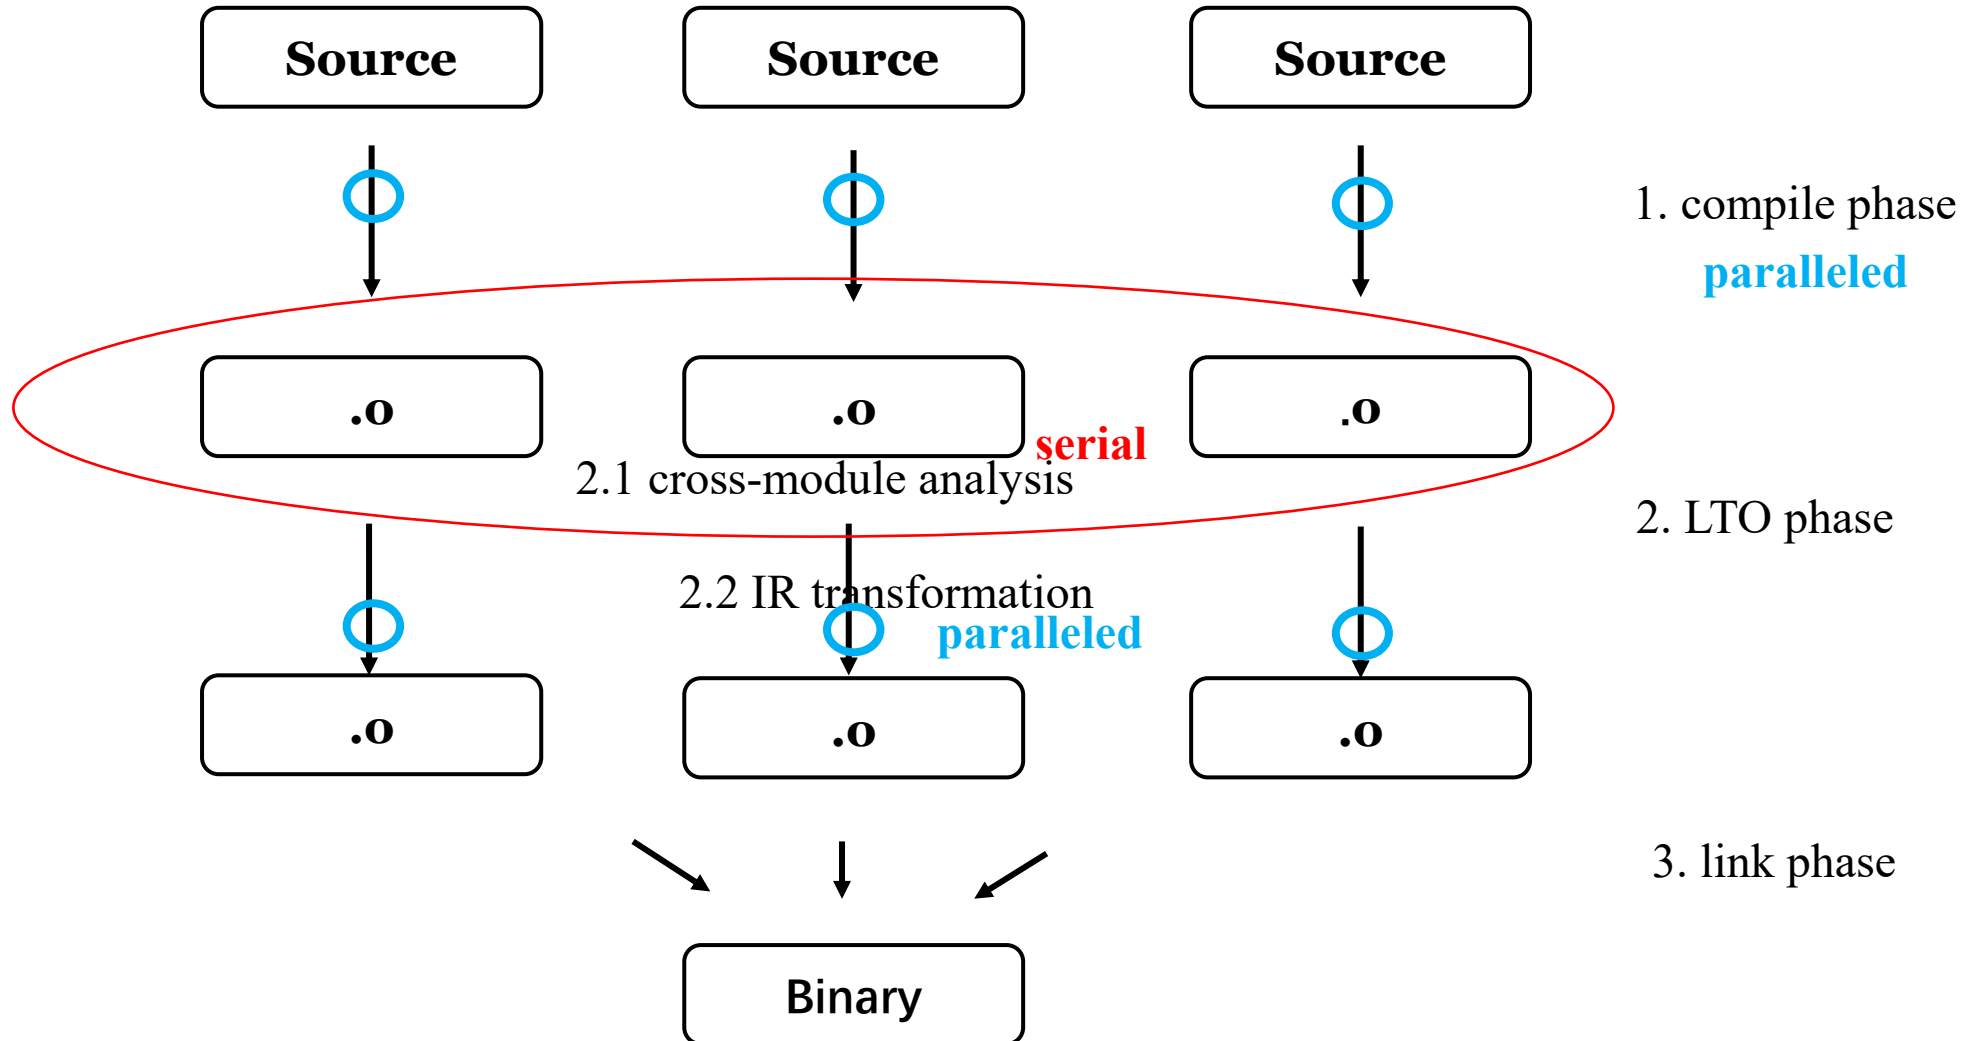
102s

2.27 GB

Previous Design of LTO

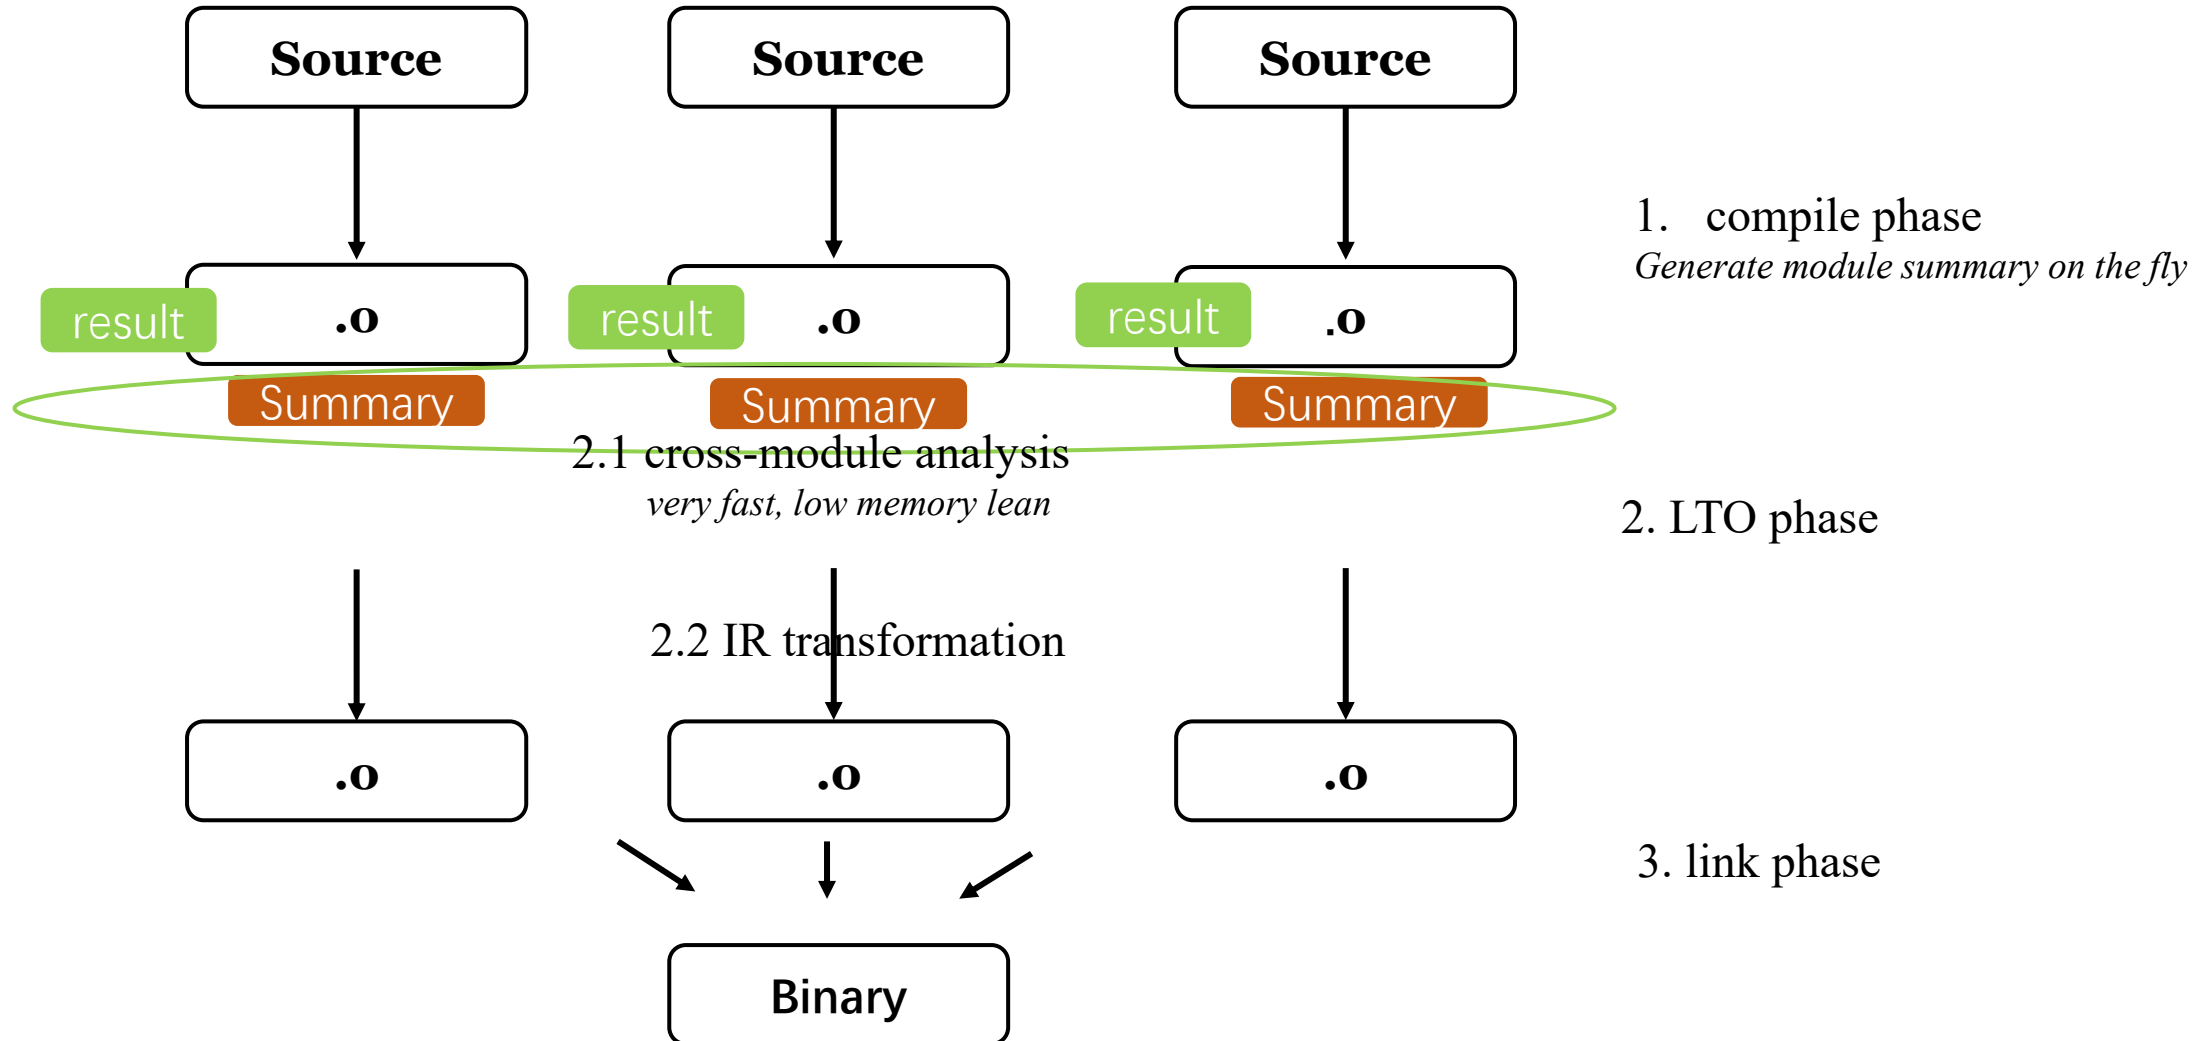
GCC, Clang(<3.9), and others

What's the
bottleneck?



ThinLTO Design

Summary-based LTO



Represent Module Summary?

A naïve inline opt under ThinLTO

main.c

```
1  int foo(int);
2  int bar(int);
3
4  int main() {
5      int x = 42;
6
7      if(...) {
8          x = bar(x);
9      } else {
10         x = foo(x);
11     }
12     return x;
13 }
```

hotness

foo() at line 10	5
bar() at line 8	30

score

foo() at line 10	20
bar() at line 8	10

A.c

```
1  int foo(int x) {
2      return 2 + x;
3  }
4
5  int bar(int y) {
6      return (x * 2) + (x * 8);
7  }
```

weigh

foo	100
bar	300

Our Inline heuristic:

A function call **f** at call site **c** has inline score:

$\text{score}(f, c) = \text{weigh}(f) / \text{hotness}(c)$

If $\text{score} < 15$ then we inline **f** at this call site.

**weigh(f) is the 100* (total operations in f).*

**hotness(c) is the called frequency recorded by profiling*

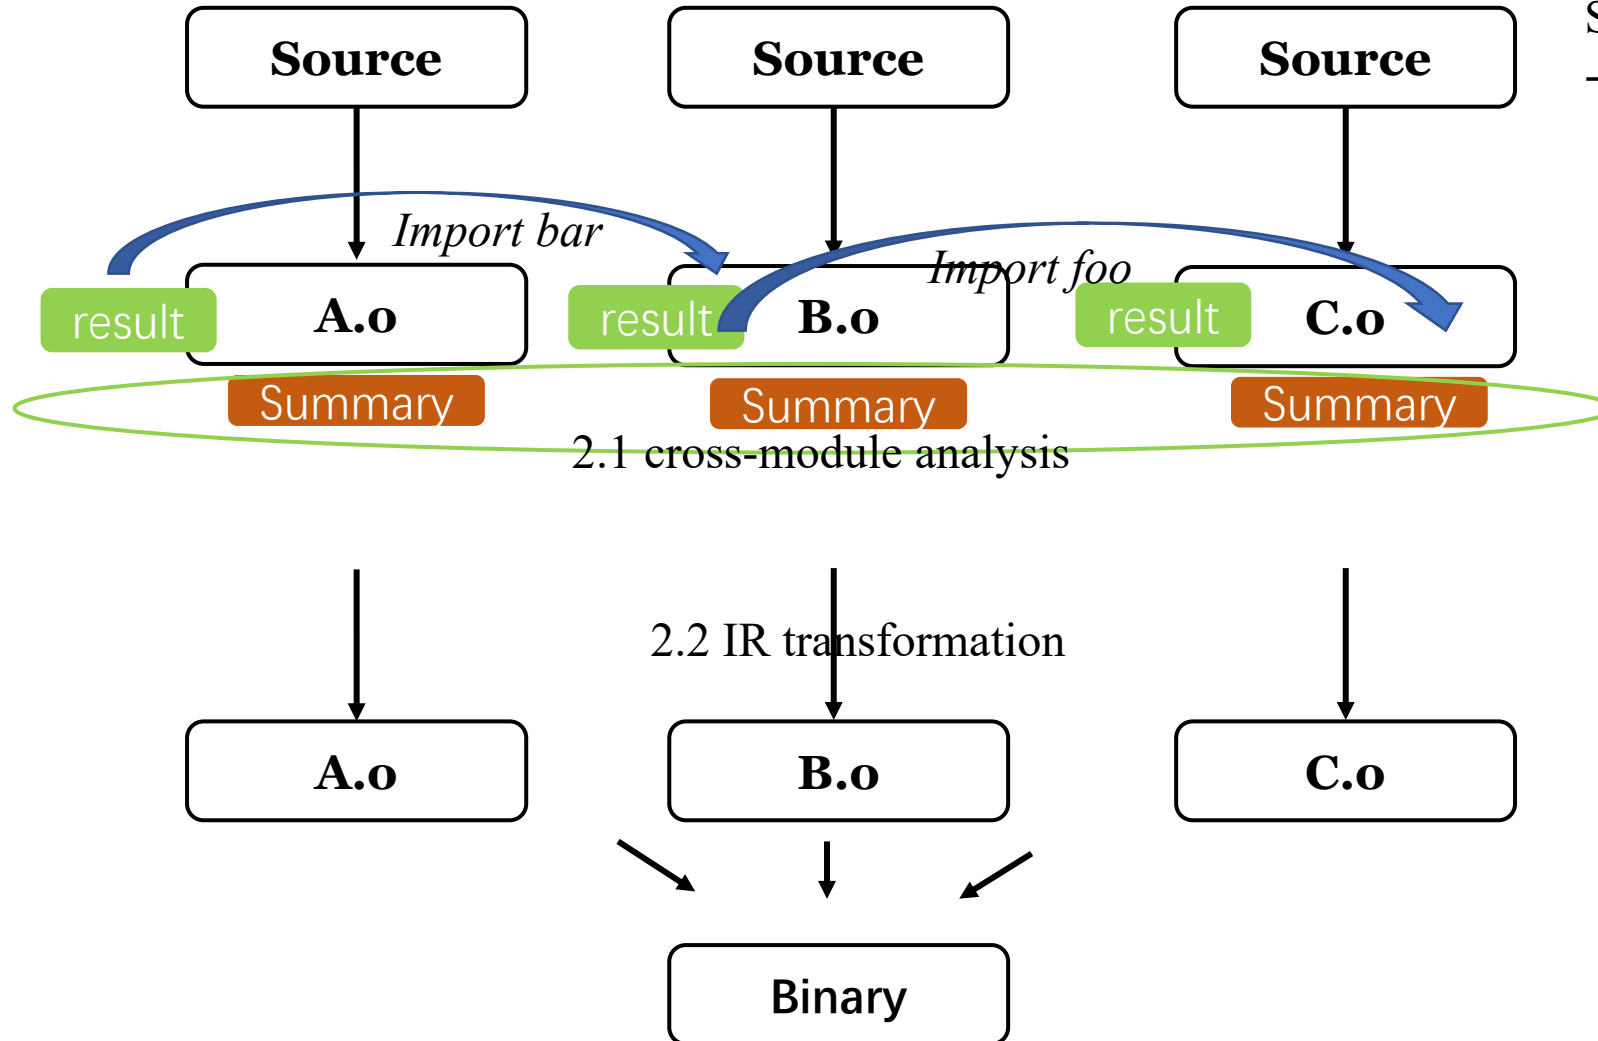
ThinLTO is a framework, not specific optimization algorithms

Design our own optimizations in ThinLTO framework

- Compile phase: collect and add information to module-summary.
 - */lib/Analysis/ModuleSummaryAnalysis.cpp*
- Cross-module analysis: gather module-summary and generate analysis result for each module.
 - *lib/LTO/ThinLTOCodeGenerator.cpp::ProcessThinLTOModule()*
- IR Transformation: For each module, perform optimizations and transformations **individually** and parallelly.
 - */lib/LTO/LTOBackend.cpp*

Distributed Build Supports

contain “import info” in analysis result



Server is assigned to do phase2.2 for **A.o**
-need **A.o + B.o**

Server is assigned to do phase2.2 for **B.o**
-need **B.o + C.o**

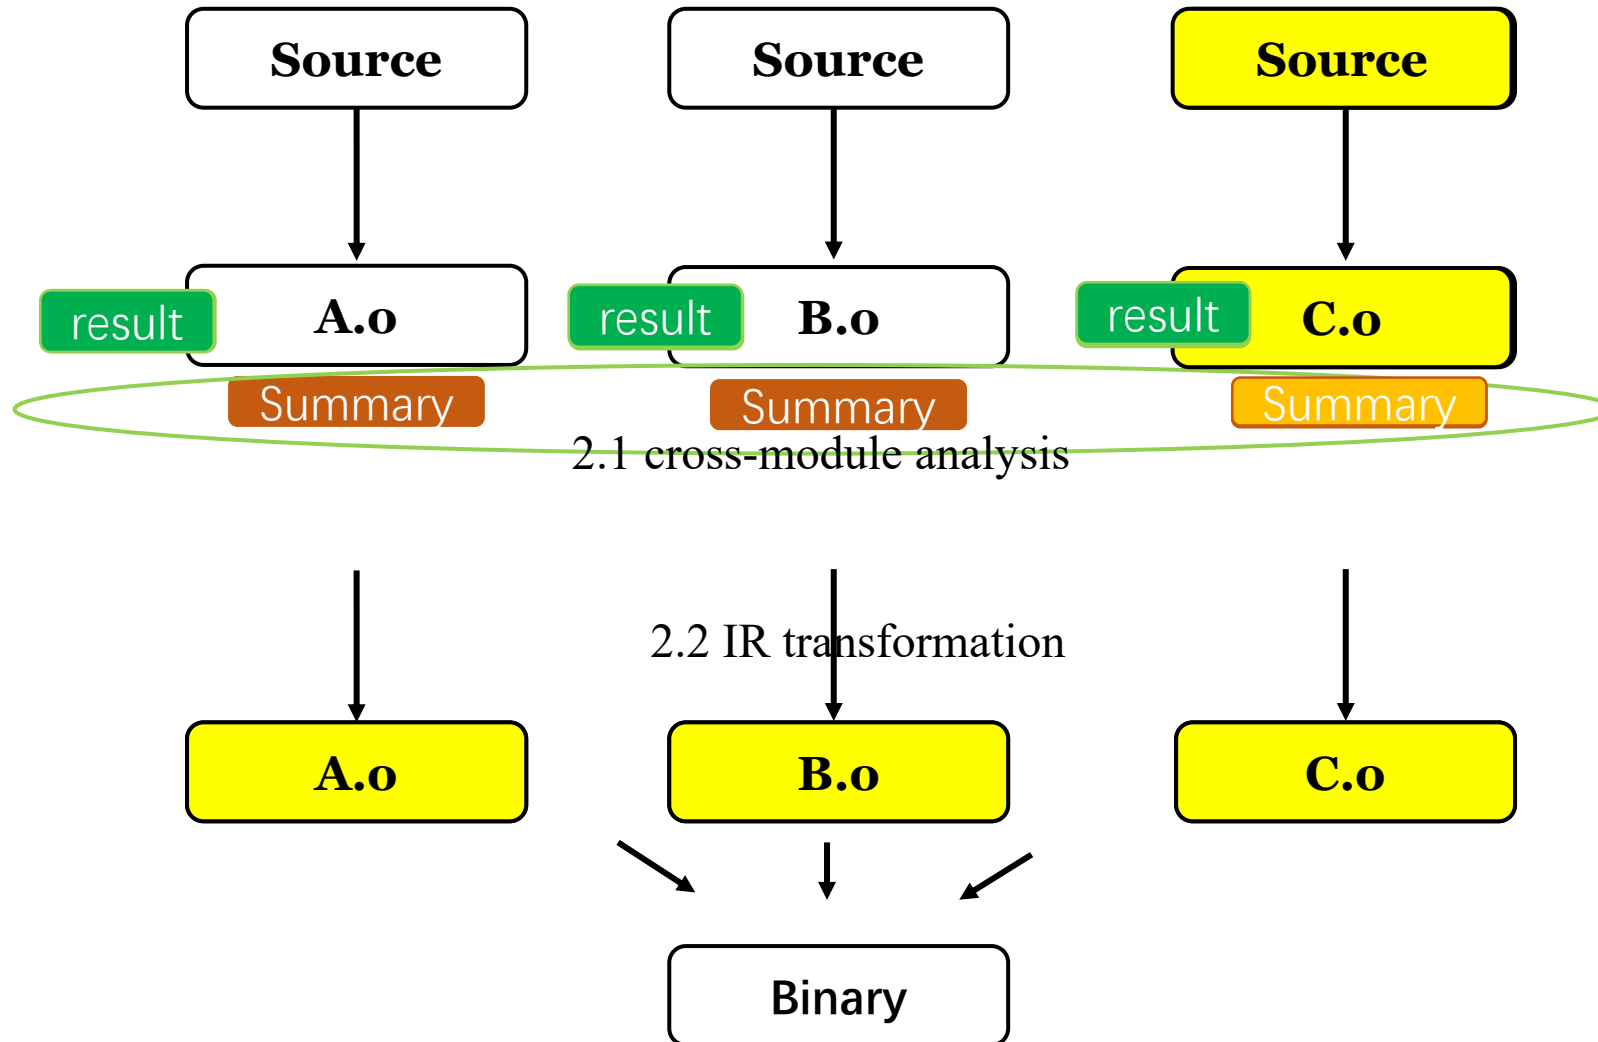
Server is assigned to do phase2.2 for **C.o**
-need **C.o**

1. compile phase

2. LTO phase

3. link phase

Incremental Build Support



1. compile phase

Naturally incremental(make)
maintain a cache system:
Source → .o file & summary

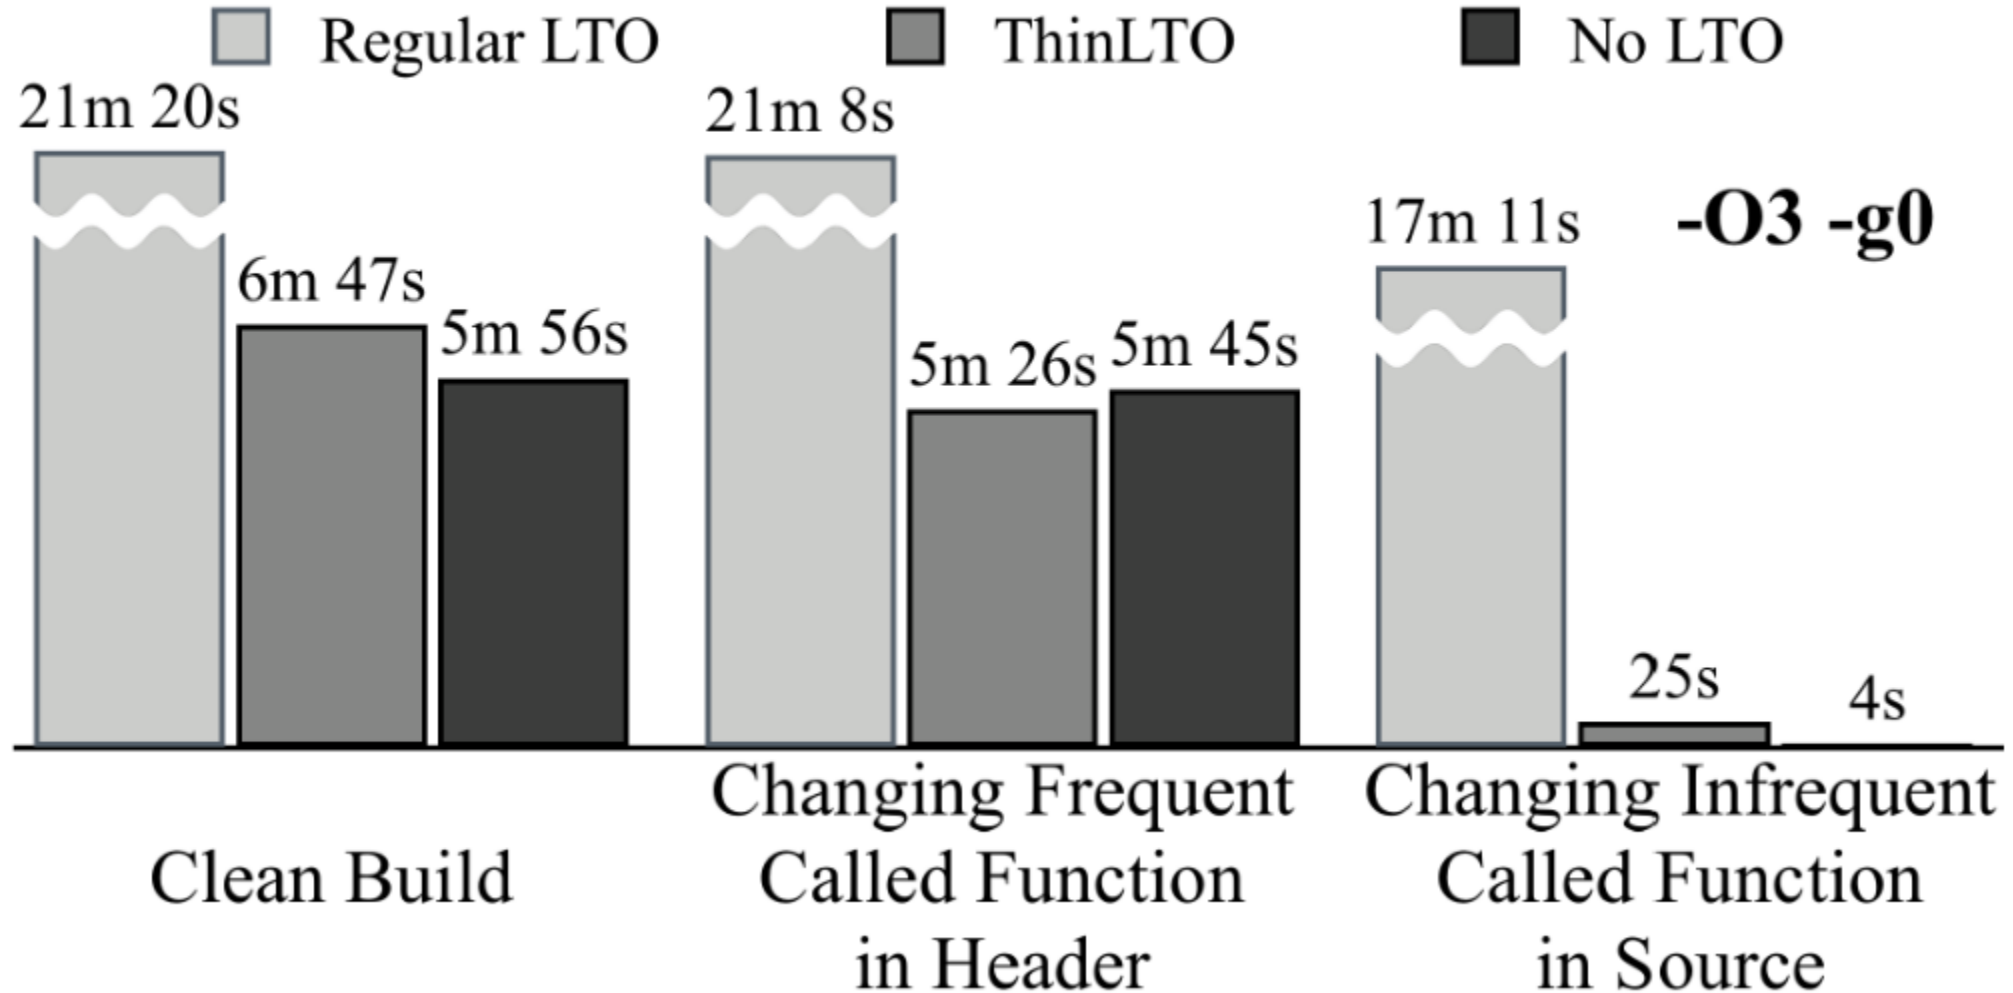
2. LTO phase

A cache system
Analysis-result → .o file

3. link phase

Incremental Build Support:

Not always so helpful?



Discussion

- Incremental build
 - Can we reduce redundant computation even when frequent called function is modified?

Discussion

Limitation: Optimize result (make program run faster)

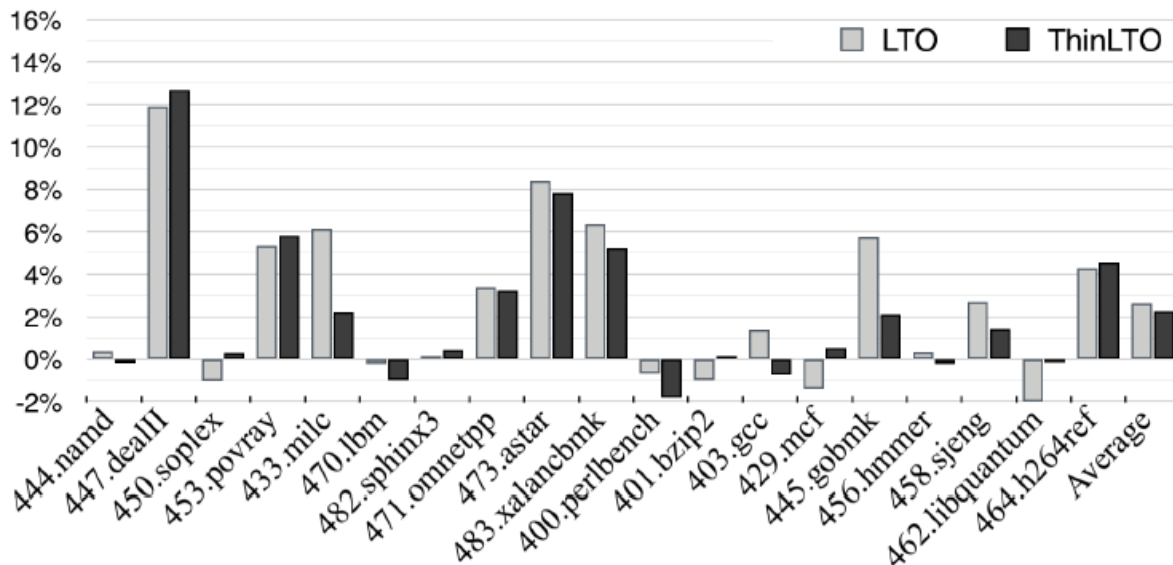
- ThinLTO's optimize result on SPEC is just “as good” as LTO.
 - ThinLTO doesn't do **aggressive** whole program analysis.

Whole program analysis

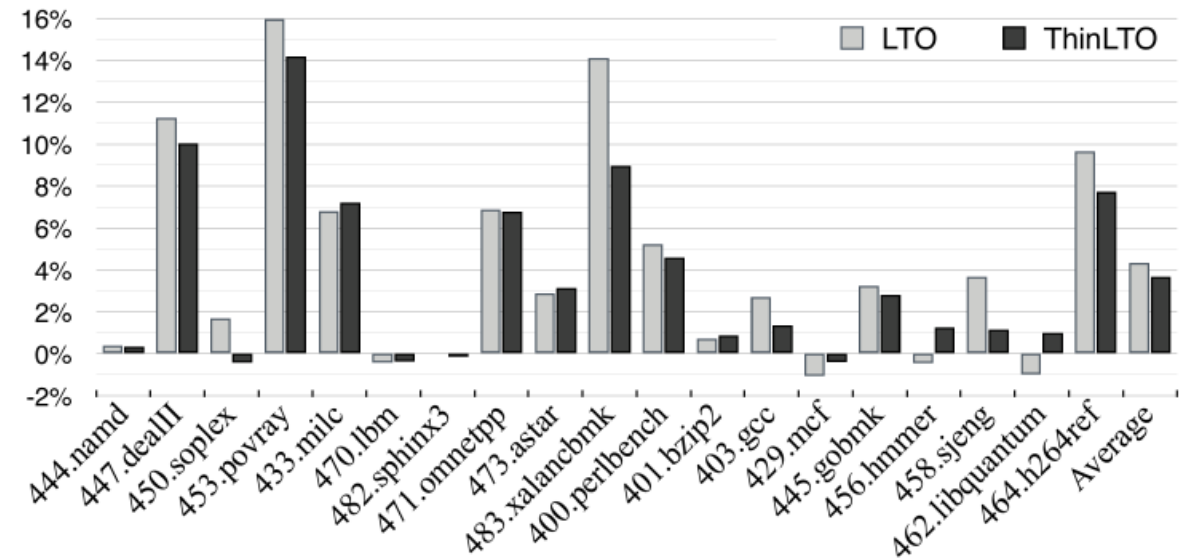


Optimize and transform for each .o file

Summary based LTO is like one-step optimization



(a) Without Profile-Guided Optimization (PGO).



(b) With Profile-Guided Optimization (PGO).

Discussion

- Optimization limitation: cont

main.c

```
1  #include <stdio.h>
2
3  int foo1(void);
4
5  void foo4(void) {
6      printf("Hi\n");
7  }
8
9
10 int main() {
11     return foo1();
12 }
```

A.c

```
1  void foo4(void);
2
3  static signed int i=0;
4
5  void foo2(void) {
6      i = -1;
7  }
8
9  static int foo3() {
10     foo4();
11     return 10;
12 }
13
14
15 int foo1(void) {
16     int data = 0;
17     if(i < 0) {
18         data = foo3();
19     }
20     data = data + 42;
21     return data;
22 }
```

Consider that **i** and **foo3** are not `static`, and are located in other files.

Can ThinLTO handle this?

Discussion

- Consider static analysis after LTO.
 - After LTO, code is smaller while semantic is the same
 - ThinLTO is fast and we can afford this cost.

Thank you!

Links

- ThinLTO: <https://dl.acm.org/doi/10.5555/3049832.3049845>
- LLVM doc: <https://clang.llvm.org/docs/ThinLTO.html>
- Talk at cpp-conf: <https://youtu.be/p9nH2vZ2mNo>
- GCC's LTO: <https://arxiv.org/abs/1010.2196>